

```
In [79]: import numpy as np
import matplotlib.pyplot as plt
import os
import pandas as pd
```

```
In [80]: df = pd.read_csv("./WA_Fn-UseC_-Telco-Customer-Churn.csv")
df.head(5)
```

Out[80]:

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines
0	7590-VHVEG	Female	0	Yes	No	1	No	No phone service
1	5575-GNVDE	Male	0	No	No	34	Yes	No
2	3668-QPYBK	Male	0	No	No	2	Yes	No
3	7795-CFOCW	Male	0	No	No	45	No	No phone service
4	9237-HQITU	Female	0	No	No	2	Yes	No

5 rows × 9 columns

notice that "TotalCharges" is object, so it may have nan values

```
In [81]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 21 columns):
customerID      7043 non-null object
gender          7043 non-null object
SeniorCitizen   7043 non-null int64
Partner         7043 non-null object
Dependents      7043 non-null object
tenure          7043 non-null int64
PhoneService    7043 non-null object
MultipleLines   7043 non-null object
InternetService 7043 non-null object
OnlineSecurity  7043 non-null object
OnlineBackup    7043 non-null object
DeviceProtection 7043 non-null object
TechSupport     7043 non-null object
StreamingTV     7043 non-null object
StreamingMovies 7043 non-null object
Contract        7043 non-null object
PaperlessBilling 7043 non-null object
PaymentMethod   7043 non-null object
MonthlyCharges  7043 non-null float64
TotalCharges    7043 non-null object
Churn           7043 non-null object
dtypes: float64(1), int64(2), object(18)
memory usage: 1.1+ MB
```

```
In [82]: df.shape
```

```
Out[82]: (7043, 21)
```

Clean up the data

```
In [83]: df['TotalCharges'] = pd.to_numeric(df['TotalCharges'], errors='coerce')
```

```
In [84]: # one-hot encoding:
# categorical variables are converted into a form that could be provided
# to ML algorithms

columns = list(df)
catg_obj = []
for attr in columns:
    if df[attr].dtype == 'object':
        if attr != 'customerID':
            catg_obj.append(attr)

df = pd.get_dummies(df, columns = catg_obj, drop_first=True)
df.head()
```

Out[84]:

	customerID	SeniorCitizen	tenure	MonthlyCharges	TotalCharges	gender_Male	Partr
0	7590-VHVEG	0	1	29.85	29.85	0	1
1	5575-GNVDE	0	34	56.95	1889.50	1	0
2	3668-QPYBK	0	2	53.85	108.15	1	0
3	7795-CFOCW	0	45	42.30	1840.75	1	0
4	9237-HQITU	0	2	70.70	151.65	0	0

5 rows × 32 columns

drop 'customerID' column

```
In [85]: df.columns
df.drop(['customerID'], axis = 1, inplace = True)
```

drop nans

```
In [86]: df.isna().sum()
```

```
Out[86]: SeniorCitizen      0
         tenure             0
         MonthlyCharges     0
         TotalCharges       11
         gender_Male        0
         Partner_Yes        0
         Dependents_Yes     0
         PhoneService_Yes   0
         MultipleLines_No phone service  0
         MultipleLines_Yes  0
         InternetService_Fiber optic    0
         InternetService_No             0
         OnlineSecurity_No internet service  0
         OnlineSecurity_Yes             0
         OnlineBackup_No internet service  0
         OnlineBackup_Yes              0
         DeviceProtection_No internet service  0
         DeviceProtection_Yes            0
         TechSupport_No internet service  0
         TechSupport_Yes                0
         StreamingTV_No internet service  0
         StreamingTV_Yes                0
         StreamingMovies_No internet service  0
         StreamingMovies_Yes            0
         Contract_One year             0
         Contract_Two year             0
         PaperlessBilling_Yes          0
         PaymentMethod_Credit card (automatic)  0
         PaymentMethod_Electronic check    0
         PaymentMethod_Mailed check       0
         Churn_Yes                    0
         dtype: int64
```

```
In [87]: df.dropna(inplace=True)
         df.shape # drop 11 rows that has nan values
```

```
Out[87]: (7032, 31)
```

```
In [88]: # df.info()
```

Train Test split

```
In [89]: from sklearn.model_selection import train_test_split

         x = df.loc[:, df.columns != 'Churn_Yes']
         y = df['Churn_Yes']
         x_train, x_test, y_train, y_test = train_test_split(x, y) # default test
         _size = 0.25
```

See if data is balance or not

```
In [90]: sum_1 = sum(y_train==1)
sum_0 = sum(y_train==0)
print sum_1
print sum_0
print (sum_1/float(sum_0+sum_1)) # Percentage churn

1415
3859
0.268297307546
```

Use SMOTE Oversampling to balance data

```
In [91]: from imblearn.over_sampling import SMOTE
smo = SMOTE()
x_train_, y_train_ = smo.fit_resample(x_train, y_train.ravel())
```

Fitting multiple models using KFold

```
In [92]: from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier

from sklearn.model_selection import cross_validate
from sklearn.model_selection import cross_val_score
from sklearn.model_selection import KFold

scoring = 'accuracy'
models = []
models.append(('LogR', LogisticRegression()))
models.append(('DT', DecisionTreeClassifier()))
models.append(('RF', RandomForestClassifier()))
rslts = []
names = []
for name, model in models:
    kf = KFold(n_splits=10, random_state=None)
    scores = cross_val_score(model, x_train_, y_train_, cv=kf, scoring=s
coring)
    rslts.append(scores)
    names.append(name)
    print "%s: %f" % (name, scores.mean())

LogR: 0.749289
DT: 0.807634
RF: 0.845593
```

Random Forest

```
In [93]: from sklearn.metrics import confusion_matrix
from sklearn.metrics import classification_report
from sklearn.metrics import accuracy_score
rf = RandomForestClassifier()
rf.fit(x_train_, y_train_)
y_pred = rf.predict(x_test)
print(accuracy_score(y_test, y_pred))

0.779294653015
```

Random Forest Confusion Matrix

```
In [94]: print(confusion_matrix(y_test, y_pred))
print(classification_report(y_test, y_pred))
```

```
[[1153  151]
 [ 237  217]]
      precision    recall  f1-score   support


      0       0.83      0.88      0.86      1304
      1       0.59      0.48      0.53       454

   micro avg       0.78      0.78      0.78      1758
   macro avg       0.71      0.68      0.69      1758
weighted avg       0.77      0.78      0.77      1758
```

Random Forest ROC Curve

```
In [95]: from sklearn import metrics

rf_pre_prob = rf.predict_proba(x_test)[: ,1]
fpr, tpr, threshold = metrics.roc_curve(y_test, rf_pre_prob)
roc_auc = metrics.auc(fpr, tpr)
plt.plot(fpr, tpr)
plt.plot([0, 1], [0, 1])
plt.savefig("RF_ROC")
```



Logistic Regression

```
In [96]: lr = LogisticRegression()
lr.fit(x_train_, y_train_)
y_pred = lr.predict(x_test)
print(accuracy_score(y_test, y_pred))
```

0.766211604096

Logistic Regression Confusion Matrix


```
In [97]: print(confusion_matrix(y_test, y_pred))
print(classification_report(y_test, y_pred))
```

```
[[977 327]
 [ 84 370]]
```

	precision	recall	f1-score	support
0	0.92	0.75	0.83	1304
1	0.53	0.81	0.64	454
micro avg	0.77	0.77	0.77	1758
macro avg	0.73	0.78	0.73	1758
weighted avg	0.82	0.77	0.78	1758

Logistic Regression ROC Curve

```
In [98]: lr_pre_prob = lr.predict_proba(x_test)[:,-1]
fpr, tpr, threshold = metrics.roc_curve(y_test, lr_pre_prob)
roc_auc = metrics.auc(fpr, tpr)
plt.plot(fpr, tpr)
plt.plot([0, 1], [0, 1])
plt.savefig("LR_ROC")
```



Using Cross validation to get a validation score

```
In [99]: from sklearn.model_selection import cross_val_score

scores = cross_val_score(rf, x_train_, y_train_, cv = 5)
print("Cross Validation Scores: " + str(scores))
print("Mean Cross Validation Score: " + str(scores.mean()))

Cross Validation Scores: [ 0.65544041  0.71567358  0.92940415  0.924222
8   0.91763943]
Mean Cross Validation Score: 0.828476072391
```

Feature importances with forests of trees



```
In [100]: features = x_train.columns
importances = rf.feature_importances_
indices = np.argsort(importances)

feature_importances = pd.DataFrame(importances, index = features, columns=
['importance']).sort_values('importance', ascending=False)
feature_importances
```

Out[100]:

	importance
TotalCharges	0.116805
tenure	0.111848
MonthlyCharges	0.101589
PaymentMethod_Electronic check	0.093270
Contract_Two year	0.076361
PaperlessBilling_Yes	0.043977
OnlineSecurity_Yes	0.042574
TechSupport_Yes	0.041188
Partner_Yes	0.039395
Contract_One year	0.038659
InternetService_Fiber optic	0.037211
Dependents_Yes	0.033305
OnlineBackup_Yes	0.027226
gender_Male	0.024512
DeviceProtection_Yes	0.023113
MultipleLines_Yes	0.019178
PaymentMethod_Credit card (automatic)	0.017200
SeniorCitizen	0.016131
StreamingMovies_Yes	0.015063
TechSupport_No internet service	0.014450
StreamingTV_Yes	0.012866
PaymentMethod_Mailed check	0.010781
StreamingMovies_No internet service	0.009302
OnlineSecurity_No internet service	0.008184
MultipleLines_No phone service	0.005781
InternetService_No	0.005288
OnlineBackup_No internet service	0.004998
PhoneService_Yes	0.004662
StreamingTV_No internet service	0.003061
DeviceProtection_No internet service	0.002024

```
In [101]: plt.figure(figsize=(15,10))
plt.title('Feature Importances')
plt.barh(range(len(indices)), importances[indices], align='center')
plt.yticks(range(len(indices)), [features[i] for i in indices])
plt.xlabel('Relative Importance')
plt.savefig("Feature Importances", dpi=150)
plt.show()
```



Top two features are TotalCharges and Tenures

```
In [102]: df.groupby(['Churn_Yes']).TotalCharges.describe()
```

Out[102]:

	count	mean	std	min	25%	50%	75%	max
Churn_Yes								
0	5163.0	2555.344141	2329.456984	18.80	577.825	1683.60	4264.125	8672.45
1	1869.0	1531.796094	1890.822994	18.85	134.500	703.55	2331.300	8684.80

```
In [103]: df.groupby(['Churn_Yes']).tenure.describe()
```

Out[103]:

	count	mean	std	min	25%	50%	75%	max
Churn_Yes								
0	5163.0	37.650010	24.076940	1.0	15.0	38.0	61.0	72.0
1	1869.0	17.979133	19.531123	1.0	2.0	10.0	29.0	72.0

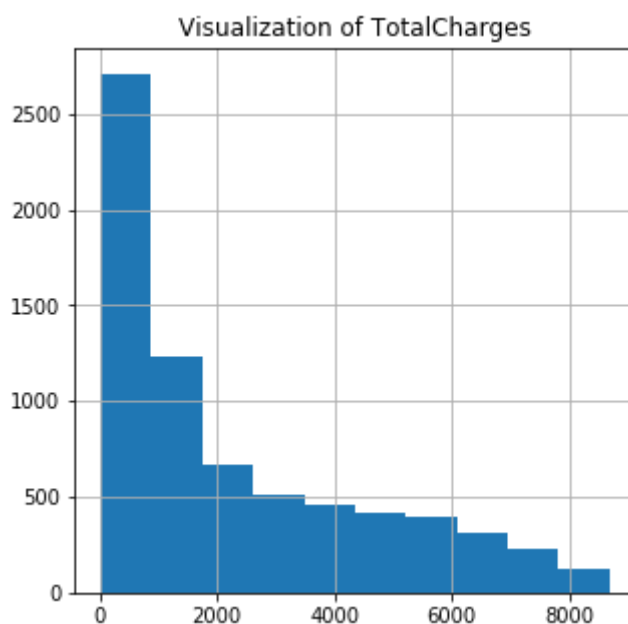
```
In [104]: num = ['float64', 'int64']  
num_data = df.select_dtypes(include=num)  
cat_data = df.select_dtypes(exclude=num)
```

```
In [105]: num_data.describe().T
```

Out[105]:

	count	mean	std	min	25%	50%	75%
SeniorCitizen	7032.0	0.162400	0.368844	0.00	0.0000	0.000	0.0000
tenure	7032.0	32.421786	24.545260	1.00	9.0000	29.000	55.0000
MonthlyCharges	7032.0	64.798208	30.085974	18.25	35.5875	70.350	89.8625
TotalCharges	7032.0	2283.300441	2266.771362	18.80	401.4500	1397.475	3794.7375

```
In [106]: df["TotalCharges"].hist(figsize=(5,5))  
plt.title('Visualization of TotalCharges')  
plt.plot  
plt.savefig("TotalCharges", dpi=300)
```



```
In [107]: df["tenure"].hist(figsize=(5,5))  
plt.title('Visualization of tenure')  
plt.plot  
plt.savefig("tenure", dpi=300)
```

